NEW YORK INSTITUTE OF TECHNOLOGY

PERSONAL PROJECT - 1

College of Engineering & Computing Science

DTSC 620 M01 - Statistics For Data Science

PROJECT ASSIGNMENT - 1

Name: AVIVARTTA KRISHNA

ID: 1261351

CLASS: DTSC 620 MO1

DATE: 11/07/2022

Professor: Dr. Kiran Balgani

Reporting Tasks:

- **Compare the accuracies of the Random Forest classifier as a function of the number of base learners (e.g., 10, 50, 100, 500, 1000, and 500) and the number of features to consider at each split (e.g., auto or sqrt). Report your observations/conclusions and provide evidence to support your conclusions. [50 points]**
- **Compare the results of all the classifiers (with the best possible parameter setting for each classifier). Use classification accuracy (# of instances correctly classified/total # of instances presented for classification), per class classification accuracy, and confusion matrix to compare the classifiers. [50 points]**

Using spam.info() on the dataset given, we are able to see that there are 58 attributes present with 55 in Float64 data type, 2 in int64 data type and one in Object Data type.

We are also able to see from the figure present below that all attributes are non-null, therefore there is no need for data cleaning.

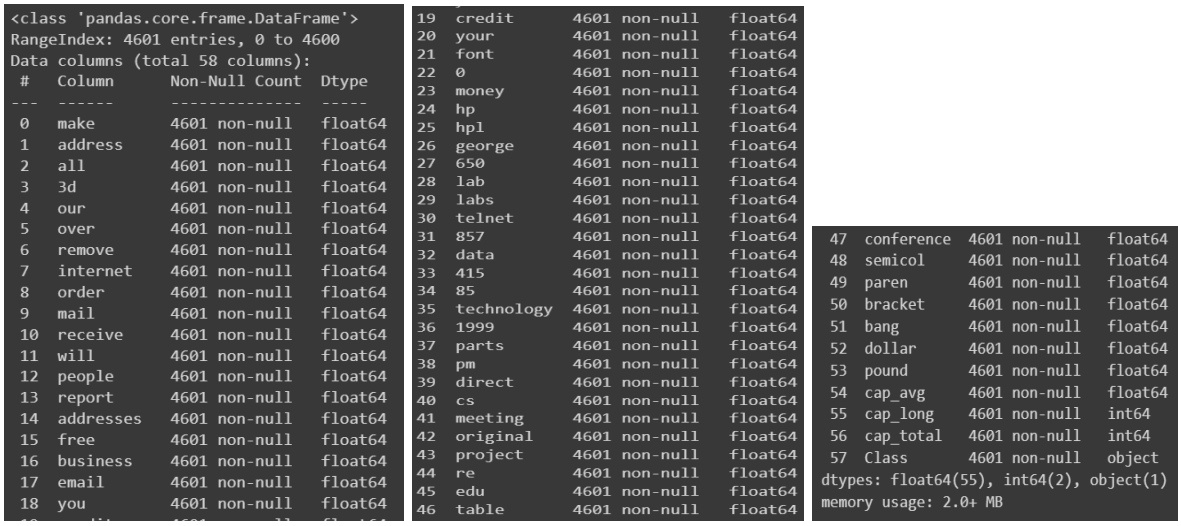

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4601 entries, 0 to 4600
Data columns (total 58 columns):
 #   Column        Non-Null Count  Dtype
---  ------        --------------  -----
 0   make          4601 non-null   float64
 1   address       4601 non-null   float64
 2   all           4601 non-null   float64
 3   3d            4601 non-null   float64
 4   our           4601 non-null   float64
 5   over          4601 non-null   float64
 6   remove        4601 non-null   float64
 7   internet      4601 non-null   float64
 8   order         4601 non-null   float64
 9   mail          4601 non-null   float64
 10  receive       4601 non-null   float64
 11  will          4601 non-null   float64
 12  people        4601 non-null   float64
 13  report        4601 non-null   float64
 14  addresses     4601 non-null   float64
 15  free          4601 non-null   float64
 16  business      4601 non-null   float64
 17  email         4601 non-null   float64
 18  you           4601 non-null   float64
 19  credit        4601 non-null   float64
 20  your          4601 non-null   float64
 21  font          4601 non-null   float64
 22  0             4601 non-null   float64
 23  money         4601 non-null   float64
 24  hp            4601 non-null   float64
 25  hpl           4601 non-null   float64
 26  george        4601 non-null   float64
 27  650           4601 non-null   float64
 28  lab           4601 non-null   float64
 29  labs          4601 non-null   float64
 30  telnet        4601 non-null   float64
 31  857           4601 non-null   float64
 32  data          4601 non-null   float64
 33  415           4601 non-null   float64
 34  85            4601 non-null   float64
 35  technology    4601 non-null   float64
 36  1999          4601 non-null   float64
 37  parts         4601 non-null   float64
 38  pm            4601 non-null   float64
 39  direct        4601 non-null   float64
 40  cs            4601 non-null   float64
 41  meeting       4601 non-null   float64
 42  original      4601 non-null   float64
 43  project       4601 non-null   float64
 44  re            4601 non-null   float64
 45  edu           4601 non-null   float64
 46  table         4601 non-null   float64
 47  conference    4601 non-null   float64
 48  semicol       4601 non-null   float64
 49  paren         4601 non-null   float64
 50  bracket       4601 non-null   float64
 51  bang          4601 non-null   float64
 52  dollar        4601 non-null   float64
 53  pound         4601 non-null   float64
 54  cap_avg       4601 non-null   float64
 55  cap_long      4601 non-null   int64
 56  cap_total     4601 non-null   int64
 57  Class         4601 non-null   object
dtypes: float64(55), int64(2), object(1)
memory usage: 2.0+ MB
```

Figure 1 - Information on each attribute present in dataset

Using print(spam.columns), we are able to see all the column names in an Array index. Visualization of the same is present below in Figure 2.

```
Index(['make', 'address', 'all', '3d', 'our', 'over', 'remove', 'internet',
       'order', 'mail', 'receive', 'will', 'people', 'report', 'addresses',
       'free', 'business', 'email', 'you', 'credit', 'your', 'font', '0',
       'money', 'hp', 'hpl', 'george', '650', 'lab', 'labs', 'telnet', '857',
       'data', '415', '85', 'technology', '1999', 'parts', 'pm', 'direct',
       'cs', 'meeting', 'original', 'project', 're', 'edu', 'table',
       'conference', 'semicol', 'paren', 'bracket', 'bang', 'dollar', 'pound',
       'cap_avg', 'cap_long', 'cap_total', 'Class'],
      dtype='object')
```

Figure 2 - Index array on all column names present in dataset.

Using spam.head(), we are able to see the first five rows of the dataset with all the information present.

	make	address	all	3d	our	over	remove	internet	order	mail	...	semicol	paren	bracket	bang	dollar	pound	cap_avg	cap_long	cap_total	Class
0	0.00	0.00	0.29	0.0	0.00	0.00	0.00	0.00	0.00	0.00	...	0.000	0.178	0.0	0.044	0.000	0.00	1.666	10	180	ham
1	0.46	0.00	0.00	0.0	0.00	0.00	0.00	0.00	0.00	0.00	...	0.000	0.125	0.0	0.000	0.000	0.00	1.510	10	74	ham
2	0.00	0.00	0.00	0.0	0.00	0.00	0.00	0.00	0.00	0.00	...	0.000	0.000	0.0	0.000	0.000	0.00	1.718	11	55	ham
3	0.33	0.44	0.37	0.0	0.14	0.11	0.00	0.07	0.97	1.16	...	0.006	0.159	0.0	0.069	0.221	0.11	3.426	72	819	spam
4	0.00	2.08	0.00	0.0	3.12	0.00	1.04	0.00	0.00	0.00	...	0.000	0.000	0.0	0.263	0.000	0.00	1.428	4	20	spam

5 rows × 58 columns

Figure 3 - first five columns of the dataset

REPORT - DECISION TREE

- In the Decision tree, we took the attribute Class as the Target Variable and every other attribute as the feature variable.
- We split the dataset into a training set and Testing Set in the ratio of 78.26% in the testing set which is approximately ~ 3601instanses and the other 1000 or 21.74% instances in the training set.
- We then created the Decision Tree Classifier. And printed the accuracy.
- Accuracy of Decision tree before optimization :

  Accuracy: 0.8850319355734518 ~ 88.5%
- Below I have given the visualization of the decision tree before optimization.

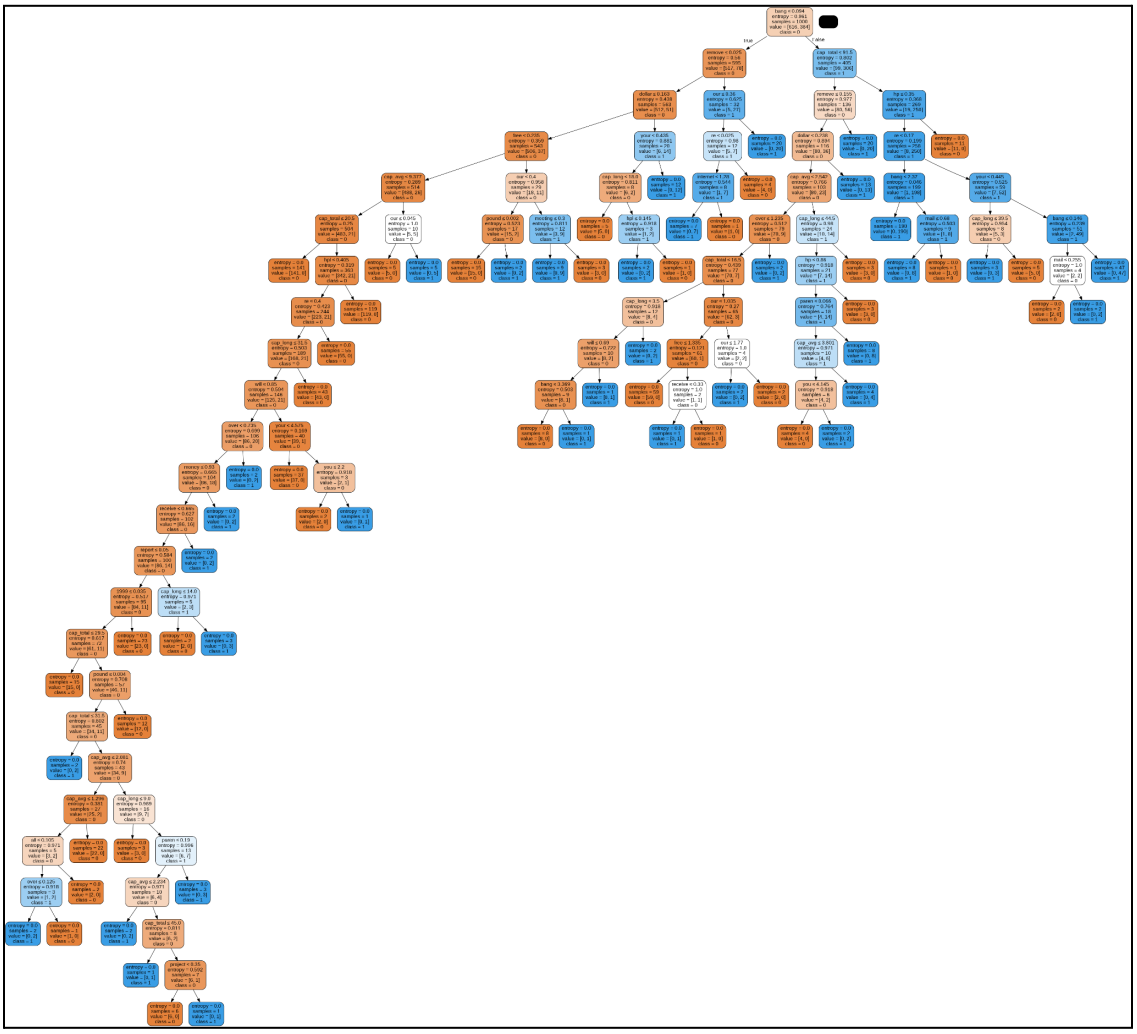

Figure 4 - Decision Tree before Optimization

Optimizing the Decision Tree

- I tried to optimize the decision tree by introducing a new criterion called entropy and trying out all the max_depth in estimation. The best optimized tree was found at max_depth = 30.
- Accuracy: 0.8733685087475701 ~ 87.3%

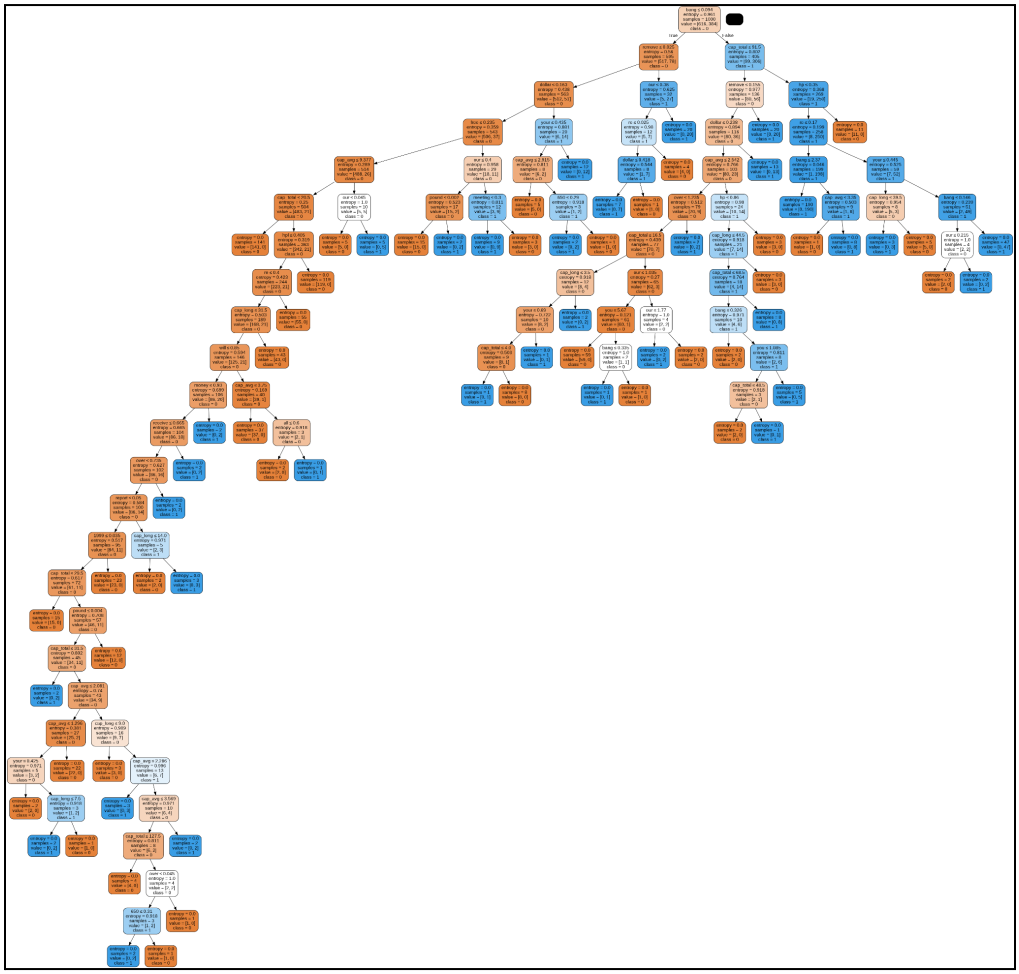

Figure 5 - Decision Tree after Optimization (very minor changes made)

Final Results for Decision Tree

DECISION TREE CLASSIFIER MODEL

	precision	recall	f1-score	support
ham	0.88	0.90	0.89	2137
spam	0.84	0.82	0.83	1464
accuracy			0.87	3601
macro avg	0.86	0.86	0.86	3601
weighted avg	0.87	0.87	0.87	3601

Figure 6 - Confusion Matrix for my Decision Tree

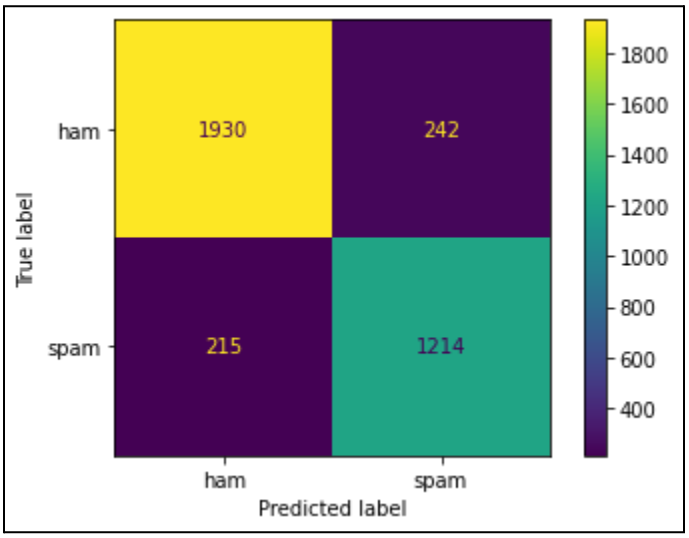

Figure 7 - Predicted Label for Spam Dataset result by Decision Tree

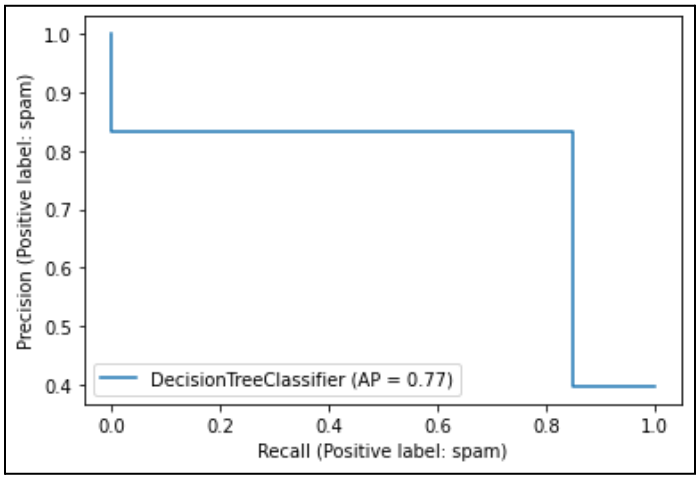

Figure 8 - Precision VS Recall of Decision Tree

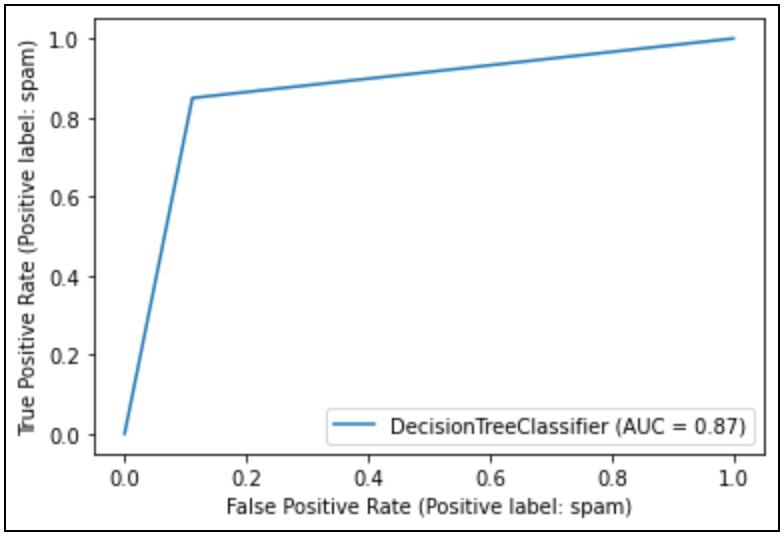

Figure 9 - True positive rate vs False positive rate for Positive Label of Spam

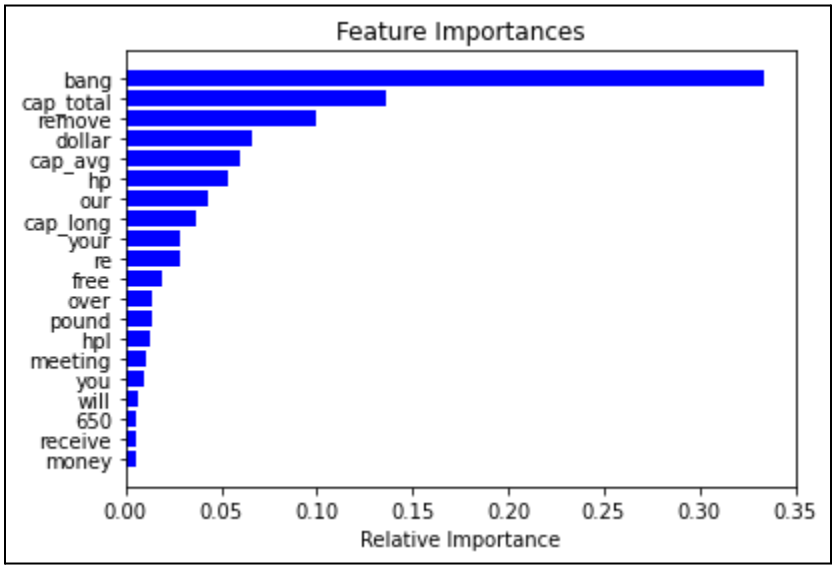

Figure 10 - Feature Importance of top 20 attributes/instances

REPORT - RANDOM FOREST CLASSIFIER

- Selecting a model with the highest accuracy given a list of base learners
  Estimators give are = {10, 50, 100, 500, 1000, 5000}
  Resultant Information:

Estimator	**Accuracy**
10	0.9255762288253263~ 92.55%
50	0.927797833935018 ~ 92.77%
100	0.9322410441544016 ~ 93.2%
500	0.9327964454318245 ~ 93.27%
1000	0.932518744793113 ~ 93.25%
5000	*0.9344626492640933 ~ 93.44%*

  Highest Accuracy = 5000

- Selecting a model with the highest accuracy given a list of Features
  Features I Chose = {None, Auto, Sqrt, Log2}
  Resultant Information

Feature	**Accuracy**
None	0.9336295473479589 ~ 93.36%
Auto	0.9327964454318245 ~ 93.27%
Sqrt	0.9305748403221328 ~ 93.05%
Log2	*0.932518744793113 ~ 93.75%*

  Highest Accuracy = Log2

RANDOM FOREST MODEL

	precision	recall	f1-score	support
ham	0.96	0.93	0.95	2253
spam	0.89	0.94	0.91	1348
accuracy			0.93	3601
macro avg	0.92	0.93	0.93	3601
weighted avg	0.93	0.93	0.93	3601

Figure 11 - Confusion Matrix for Random Forest

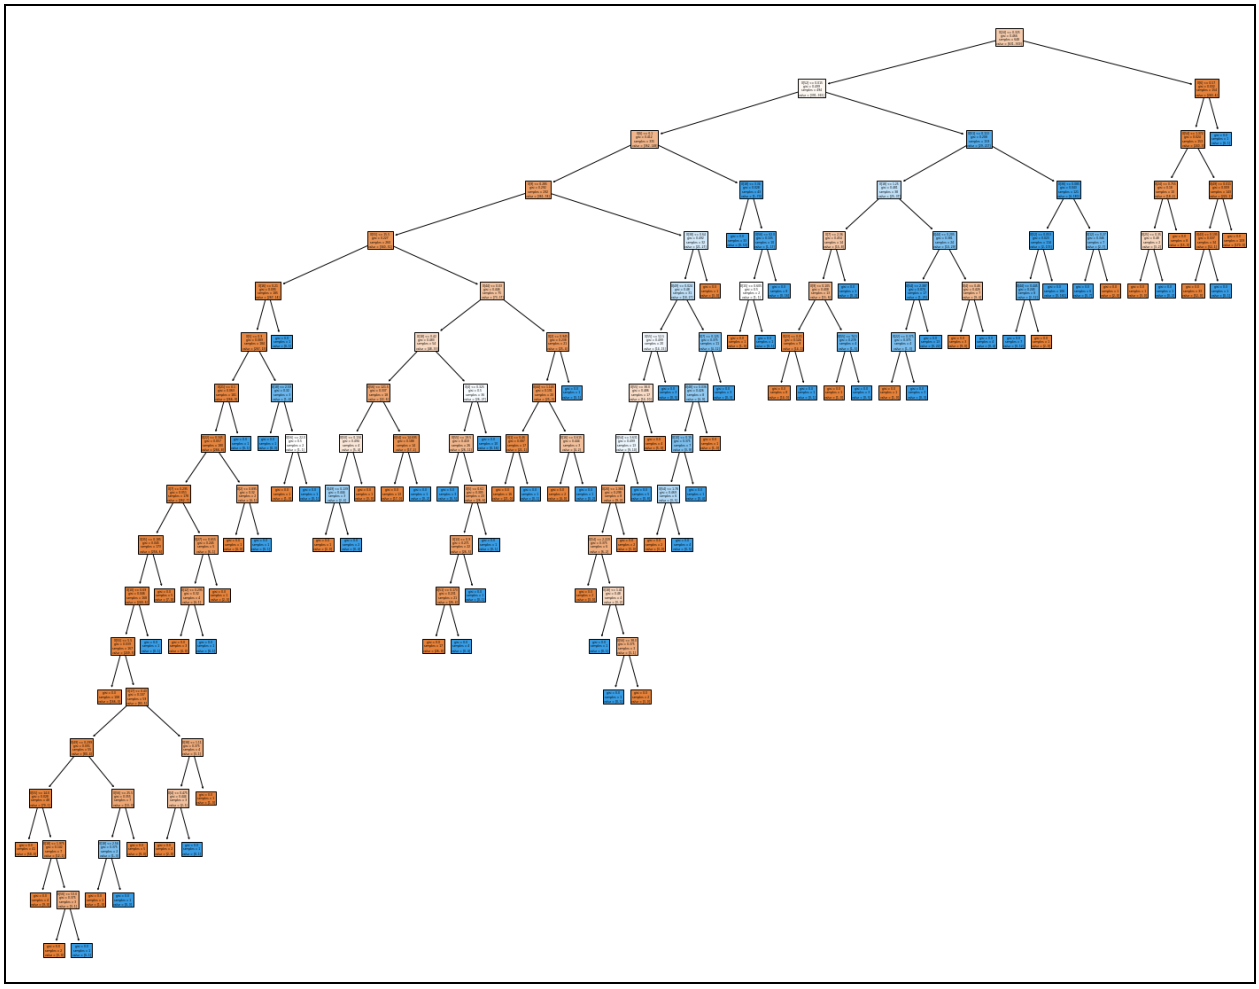

Figure 12 - Random Forest Classifier Visualization

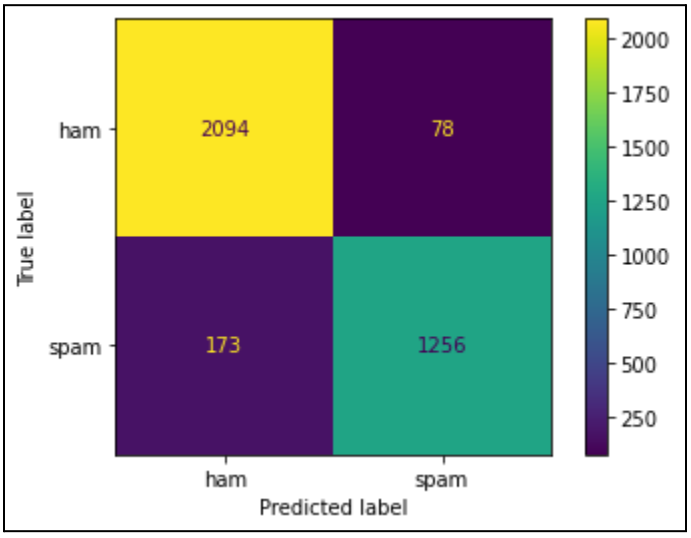

Figure 13 - Predicted Label for Spam Dataset result by Random Forest Classifier

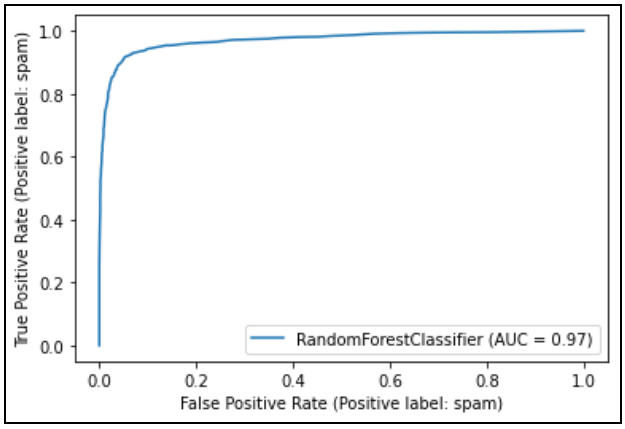

Figure 14 - True Positive VS False Positive for Positive Label of Spam

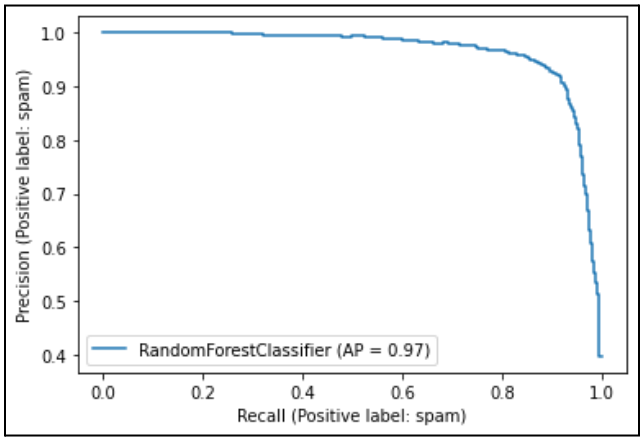

Figure 15 - Precision VS Recall for Positive Label - Spam

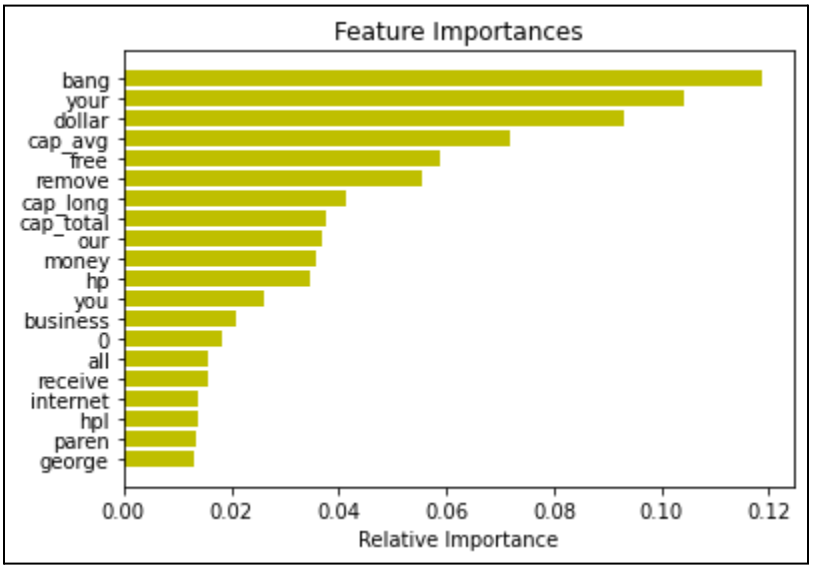

Figure 10 - Feature Importance of top 20 attributes/instances

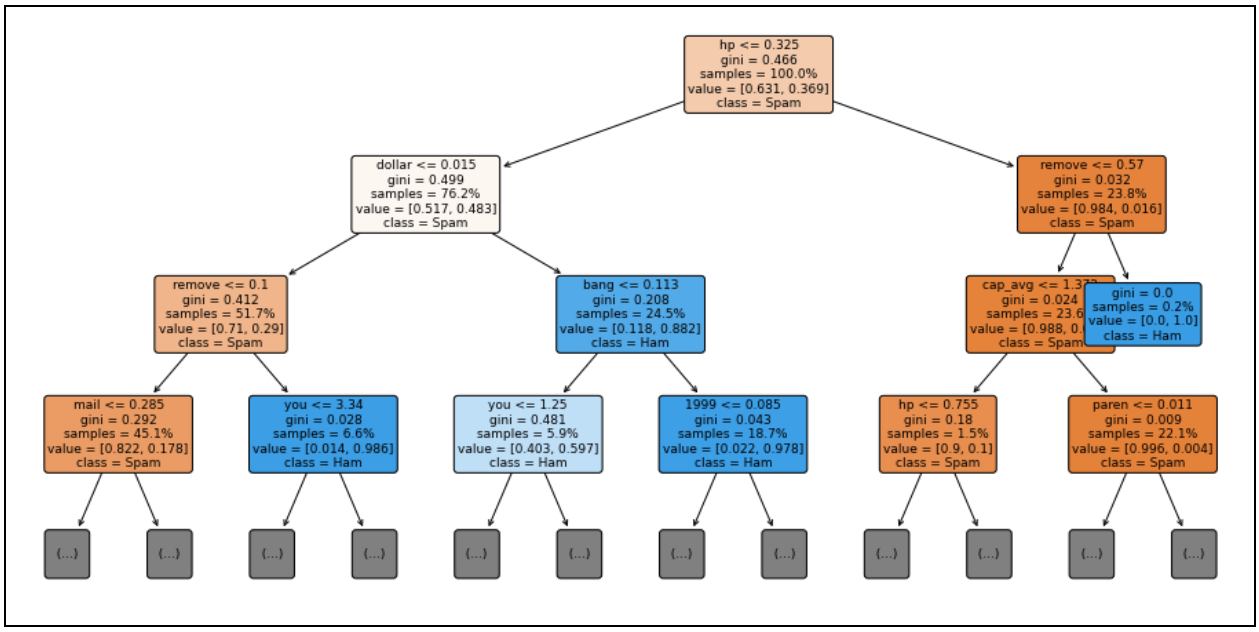

Figure 11 - Optimized Random Forest

FINAL STATEMENT : The Random Forest Classifier has a better performance since it classified 2094 data points correctly as opposed to the decision tree which only classified 1930 data points correctly.

Link to Google Colab CODE:
https://colab.research.google.com/drive/1g0_n5_YsMR2tBWrCS7ZTC_N_fZ2JKLwH
